DAY 10 OF 30

Dataflows Gen2

Cleaning data without code.
Where it shines, and where it stops.

Power Query | Low-code | Destinations

Keshav Maheshwari
Data Engineer | Databricks · AWS & Azure · Streaming · Applied AI

Swipe →

THE PROBLEM

Not everyone who cleans data writes code.

1	**The analyst**	Knows the data best
2	**The tool they use**	Excel, by hand, every month
3	**The risk**	Steps live in one person's head

The logic is real and valuable. **It is just not written down anywhere a machine can run it.**

Keshav Maheshwari
Data Engineer | Databricks · AWS & Azure · Streaming · Applied AI

Swipe →

THE SIMPLE VERSION

A recorded recipe.

You clean the data once by clicking: remove this column, split that one, filter those rows. Every step is recorded.

From then on the recipe **runs by itself**, the same way every time.

✓ **Built with**	Power Query, in the browser
✓ **Steps**	Recorded and editable
✓ **Output**	Lands in a table you choose

Keshav Maheshwari
Data Engineer | Databricks · AWS & Azure · Streaming · Applied AI

Swipe →

HOW IT WORKS

How a dataflow runs

1. **Connect** to a source from a long list of connectors
2. **Shape** the data with point-and-click steps
3. **Choose a destination**: Lakehouse, Warehouse and more
4. **Refresh** on a schedule or from a pipeline

Same engine as Excel and Power BI

Anyone who has used Power Query already knows how to build one.

Keshav Maheshwari
Data Engineer | Databricks · AWS & Azure · Streaming · Applied AI

Swipe →

FOR ENGINEERS

What engineers should know

Query folding decides the speed

When steps can be pushed down to the source, it is fast. When they cannot, rows are pulled and processed in the engine.

Staging has a cost

Staging helps large transformations but adds compute. Turn it off for simple loads.

Fast copy for big moves

Large, simple ingestion can skip row-by-row processing.

Git and CI/CD are built in

Since April 2026 every new dataflow is created with Git and deployment support, like other code.

Keshav Maheshwari
Data Engineer | Databricks · AWS & Azure · Streaming · Applied AI

Swipe →

THE HONEST PART

Dataflow or notebook?

DATAFLOW

Small to medium data

Analysts own the logic

Many sources with simple shaping

NOTEBOOK

Very large tables

Complex logic that needs tests

Tight control over performance

> **The catch:** a dataflow that was fine at one million rows can struggle at one hundred million. Watch the refresh time, and move heavy steps to Spark before they hurt.

Keshav Maheshwari
Data Engineer | Databricks · AWS & Azure · Streaming · Applied AI

Swipe →

WATCH OUT

Four things to watch

1. **Steps that break folding** early in the query
2. **Staging left on** for loads that do not need it
3. **One giant dataflow** instead of small, reusable ones
4. **No owner** once the analyst who built it moves on

Low-code is still code. **It needs the same care.**

Keshav Maheshwari
Data Engineer | Databricks · AWS & Azure · Streaming · Applied AI

Swipe →

SAVE THIS

Dataflows in five lines

1. **Power Query** in the browser
2. **Every step** is recorded and repeatable
3. **Writes** to a Lakehouse or Warehouse
4. **Great** for analysts and modest volumes
5. **Spark** takes over when data gets big

Tomorrow: Notebooks and Spark. Where the heavy processing happens.

Keshav Maheshwari
Data Engineer | Databricks · AWS & Azure · Streaming · Applied AI

Follow for Day 11